

First, reproduce learning in a simpler setting.

A recursive language model (RLM) can use code and ask helper models to solve parts of a problem. We eventually want to train those choices.

Today's question

Can we use a published training recipe to improve a model's math answers on questions excluded from training?

We use the authors' code for **Dr. GRPO**, a reinforcement learning (RL) method. It learns from scores assigned to the model's own answers.

First check: The authors' released model scored 73.2% in our test, versus 74.2% reported. That checks their saved model, not our own training.

Liu et al., *Understanding R1-Zero-Like Training: A Critical Perspective* (2025).

Official code; our run is much smaller than the paper's experiments.

Reward correct attempts, then change the model.

Illustrative question: Solve $3x + 2 = 14$. The correct answer is $x = 4$.

The model's attempts	Checker score	Learning direction
Two answers end in $x = 4$.	1 each	Make these more likely.
Two end in $x = 5$ or $x = 6$.	0 each	Make these less likely.

Each attempt is compared with the average score for that question. These differences guide changes to the model's numerical settings, or **weights**.

The checker judges the final answer. We do not supply worked solutions for the model to copy.

Simplified illustration with four attempts; our runs use eight. If every attempt receives the same score, this question supplies no relative learning signal.

We train on math, then test on different questions.

We add RL to the pretrained math model **Qwen2.5-Math-1.5B**.

Training budget	Questions	Weight updates
Completed short runs	512	32
Longer run now underway	4,080	255 planned

We test on **MATH500: 500 math questions excluded from our weight training**. A fixed subset of 64 tracks progress.

Two ways of asking the same question:

Chat style: A conversation with instructions to explain the steps and mark the final answer.

Question alone: Give only the original question.

Every run starts from the original model; one A100 GPU, eight answers per training question. The 64 monitored questions are not an untouched final test.

The input format changes the size of the gain.

Correct answers out of the same 500 MATH500 questions.


Model tested	Chat-style input	Question alone
Starting model	154	305
Trained with chat style	308	317
Trained with questions alone	168	314
Repeated question-only training	166	306

Question-only training scored 314 and 306, versus 305 before training. We have not established a reliable gain over that stronger starting setup.

Each trained row is a separate run at its final checkpoint. The same checker and length limit apply throughout. Scores judge final answers, not every reasoning step.

Longer training: an early gain, then a level score.

Correct final answers out of the same 64 monitoring questions.



One training run, with the same questions, chat-style input, checker, and answer-length limit.

At the last check, five answers improved and five became wrong. The total stayed at 42.

Next: Finish the fixed training budget, then test all 500 questions and two other math collections: AMC and Minerva.

Interim result from one ongoing run, not a final benchmark score. We have not yet shown that this longer recipe beats the stronger starting setup or the short runs.